

---

# COMP 551 – Applied Machine Learning

## Lecture 2: Linear regression

---

**Instructor:** Joelle Pineau ([jpineau@cs.mcgill.ca](mailto:jpineau@cs.mcgill.ca))

**Class web page:** [www.cs.mcgill.ca/~jpineau/comp551](http://www.cs.mcgill.ca/~jpineau/comp551)

Unless otherwise noted, all material posted for this course are copyright of the instructor, and cannot be reused or reposted without the instructor's written permission.

---

---

## Today's Quiz (informal)

---

Write down the 3 most useful insights you gathered from the article:

*"A Few Useful Things to Know About Machine Learning".*

- domain-specific  $\equiv$  feature selection
- no free lunch.
- simplicity  $\neq$  usefulness

# Supervised learning

- Given a set of **training examples**:  $x_i = \langle x_{i1}, x_{i2}, x_{i3}, \dots, x_{in}, y_i \rangle$   
 $x_{ij}$  is the  $j^{th}$  feature of the  $i^{th}$  example  
 $y_i$  is the desired **output** (or **target**) for the  $i^{th}$  example.  
 $X_j$  denotes the  $j^{th}$  feature.
- We want to learn a function  $f: X_1 \times X_2 \times \dots \times X_n \rightarrow Y$   
which maps the input variables onto the output domain.

tumor size	texture	perimeter	...	outcome	time
18.02	27.6	117.5		N	31
17.99	10.38	122.8		N	61
20.29	14.34	135.1		R	27
...					

---

# Supervised learning

---

- Given a dataset  $X \times Y$ , find a function:  $f: X \rightarrow Y$  such that  $f(\mathbf{x})$  is a good predictor for the value of  $y$ .
- Formally,  $f$  is called the **hypothesis**.
- Output  $Y$  can have many types:
  - If  $Y = \mathbb{R}$ , this problem is called **regression**.
  - If  $Y$  is a finite discrete set, the problem is called **classification**.
  - If  $Y$  has 2 elements, the problem is called **binary classification**.

---

# Prediction problems

---

- The problem of predicting tumour recurrence is called:

**classification**

- The problem of predicting the time of recurrence is called:

**regression**

- Treat them as two separate supervised learning problems.

tumor size	texture	perimeter	...	outcome	time
18.02	27.6	117.5		N	31
17.99	10.38	122.8		N	61
20.29	14.34	135.1		R	27
...					

---

# Variable types

---

- **Quantitative**, often real number measurements.
  - Assumes that similar measurements are similar in nature.
- **Qualitative**, from a set (categorical, discrete).
  - E.g. {Spam, Not-spam}
- **Ordinal**, also from a discrete set, without metric relation, but that allows ranking.
  - E.g. {first, second, third}

---

# The i.i.d. assumption

---

- In supervised learning, the examples  $x_i$  in the training set are assumed to be independently and identically distributed.
  - **Independently**: Every  $x_i$  is freshly sampled according to some probability distribution  $D$  over the data domain  $X$ .
  - **Identically**: The distribution  $D$  is the same for all examples.
- **Why?**

---

# Empirical risk minimization

---

For a given function class  $F$  and training sample  $S$ ,

- Define a notion of error (*left intentionally vague for now*):

$$L_S(f) = \# \text{ mistakes made on the sample } S$$

- Define the Empirical Risk Minimization (ERM):

$$ERM_F(S) = \operatorname{argmin}_{f \in F} L_S(f)$$


where *argmin* returns the function  $f$  (or set of functions) that achieves the minimum loss on the training sample.

- Easier to minimize the error with i.i.d. assumption.



# A regression problem

- What hypothesis class should we pick?



Observe	Predict
<u>x</u>	<u>y</u>
0.86	2.49
0.09	0.83
-0.85	-0.25
0.87	3.10
-0.44	0.87
-0.43	0.02
-1.1	-0.12
0.40	1.81
-0.96	-0.83
0.17	0.43

---

# Linear hypothesis

---

- Suppose  $Y$  is a **linear function** of  $X$ :

$$\begin{aligned}f_w(\mathbf{X}) &= w_0 + w_1 x_1 + \dots + w_m x_m \\ &= w_0 + \sum_{j=1:m} w_j x_j\end{aligned}$$

- The  $w_j$  are called **parameters** or **weights**.
- To simplify notation, we add an attribute  $x_0=1$  to the  $m$  other attributes (also called **bias term** or **intercept**).

**How should we pick the weights?**

# Least-squares solution method

- The linear regression problem:  $f_w(X) = w_0 + \sum_{j=1:m} w_j x_j$   
where  $m$  = the dimension of observation space, i.e. number of features.

- Goal:** Find the **best** linear model given the data.
- Many different possible **evaluation** criteria!
- Most common choice is to find the  $w$  that minimizes:

$$Err(w) = \sum_{i=1:n} (y_i - w^T x_i)^2$$


(A note on notation: Here  $w$  and  $x$  are column vectors of size  $m+1$ .)

$$Err(w) = \sum_{i=1:n} (y_i - w^T x_i)^2$$

---

# Least-squares solution for $X \in \mathbb{R}^2$

---



---

# Least-squares solution method

---

- Re-write in matrix notation:  $f_w(X) = Xw$

$$Err(w) = (Y - Xw)^T (Y - Xw)$$

where  $X$  is the  $n \times m$  matrix of input data,  
 $Y$  is the  $n \times 1$  vector of output data,  
 $w$  is the  $m \times 1$  vector of weights.

- To minimize, take the derivative w.r.t.  $w$ :

$$\partial Err(w) / \partial w = -2 X^T (Y - Xw)$$

- You get a system of  $m$  equations with  $m$  unknowns.

- Set these equations to 0:  $X^T (Y - Xw) = 0$


---

# Least-squares solution method

---

- We want to solve for  $\mathbf{w}$ :  $X^T (Y - X\mathbf{w}) = 0$
- Try a little algebra:  $X^T Y = X^T X \mathbf{w}$   
 $\hat{\mathbf{w}} = (X^T X)^{-1} X^T Y$   
 $\hat{\mathbf{w}} = (X^T X)^{-1} X^T Y$   
( $\hat{\mathbf{w}}$  denotes the estimated weights)
- The fitted data:  $\hat{Y} = X\hat{\mathbf{w}} = X (X^T X)^{-1} X^T Y$
- To predict new data  $X' \rightarrow Y'$ :  $Y' = X'\hat{\mathbf{w}} = X' (X^T X)^{-1} X^T Y$

# Example of linear regression



$$w = (X^T X)^{-1} X^T Y$$

$x$	$y$
0.86	2.49
0.09	0.83
-0.85	-0.25
0.87	3.10
-0.44	0.87
-0.43	0.02
-1.10	-0.12
0.40	1.81
-0.96	-0.83
0.17	0.43

What is a plausible estimate of  $w$  ?

**Try it!**

# Data matrices

$$\begin{aligned}
 X^T X &= \begin{bmatrix} 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 0.86 & 1 \\ 0.09 & 1 \\ -0.85 & 1 \\ 0.87 & 1 \\ -0.44 & 1 \\ -0.43 & 1 \\ -1.10 & 1 \\ 0.40 & 1 \\ -0.96 & 1 \\ 0.17 & 1 \end{bmatrix} \\
 &= \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix}
 \end{aligned}$$

$n \times m$   $m \times 1$   
 $m^2$



---

# Data matrices

---


$$X^T Y =$$
$$\begin{bmatrix} 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 2.49 \\ 0.83 \\ -0.25 \\ 3.10 \\ 0.87 \\ 0.02 \\ -0.12 \\ 1.81 \\ -0.83 \\ 0.43 \end{bmatrix}$$
$$= \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix}$$

# Solving the problem

$$\mathbf{w} = (X^T X)^{-1} X^T Y = \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 1.60 \\ 1.05 \end{bmatrix}$$

$\beta_1$   
 $\beta_0$

So the best fit line is  $y = 1.60x + 1.05$ .



---

# Interpreting the solution

---

- Linear fit for a prostate cancer dataset
  - Features  $X = \{\text{lcavol}, \text{lweight}, \text{age}, \text{lbph}, \text{svi}, \text{lcp}, \text{gleason}, \text{pgg45}\}$
  - Output  $y$  = level of PSA (an enzyme which is elevated with cancer).
  - High coefficient weight (in absolute value) = important for prediction.

Term	Coefficient	Std. Error
Intercept	$w_0 = 2.46$	0.09
lcavol	0.68	0.13
lweight	0.26	0.10
age	-0.14	0.10
lbph	0.21	0.10
svi	0.31	0.12
lcp	-0.29	0.15
gleason	-0.02	0.15
pgg45	0.27	0.15

# Computational cost of linear regression

- What operations are necessary?  $\sim nm^2 + m^3$ 
  - Overall: 1 matrix inversion + 3 matrix multiplications
  - $X^T X$  (other matrix multiplications require fewer operations.)
    - $X^T$  is  $m \times n$  and  $X$  is  $n \times m$ , so we need  $nm^2$  operations.
  - $(X^T X)^{-1}$ 
    - $X^T X$  is  $m \times m$ , so we need  $m^3$  operations.
- We can do linear regression in polynomial time, but handling large datasets (many examples, many features) can be problematic.

# Gradient Solution

An alternative for minimizing mean-squared error (MSE)


- Recall the least-square solution:  $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$
- What if  $\mathbf{X}$  is too big to compute this explicitly (e.g.  $m \sim 10^6$ )?
- Go back to the gradient step:  $Err(\mathbf{w}) = (\mathbf{Y} - \mathbf{X}\mathbf{w})^T (\mathbf{Y} - \mathbf{X}\mathbf{w})$

$$\partial Err(\mathbf{w}) / \partial \mathbf{w} = -2 \mathbf{X}^T (\mathbf{Y} - \mathbf{X}\mathbf{w})$$

$$\partial Err(\mathbf{w}) / \partial \mathbf{w} = 2(\mathbf{X}^T \mathbf{X}\mathbf{w} - \mathbf{X}^T \mathbf{Y})$$

# Gradient-descent solution for MSE

- Consider the error function:



- The gradient of the error is a vector indicating the direction to the minimum point.
- Instead of directly finding that minimum (using the closed-form equation), we can take small steps towards the minimum.

# Gradient-descent solution for MSE

- We want to produce a sequence of weight solutions,  $\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots$ , such that:  $Err(\mathbf{w}_0) > Err(\mathbf{w}_1) > Err(\mathbf{w}_2) > \dots$

- The algorithm:

*Given an initial weight vector  $\mathbf{w}_0$ ,*

*Do for  $k=1, 2, \dots$*

$$\mathbf{w}_{k+1} = \mathbf{w}_k - \alpha_k \partial Err(\mathbf{w}_k) / \partial \mathbf{w}_k$$

*End when  $|\mathbf{w}_{k+1} - \mathbf{w}_k| < \epsilon$*

- Parameter  $\alpha_k > 0$  is the step-size (or learning rate) for iteration  $k$ .

---

# Convergence

---

- Convergence depends in part on the  $\alpha_k$ .
- **If steps are too large**: the  $w_k$  may oscillate forever.
  - This suggests that  $\alpha_k \rightarrow 0$  as  $k \rightarrow \infty$ .
- **If steps are too small**: the  $w_k$  may not move far enough to reach a local minimum.



---

# Robbins-Monroe conditions

---

- The  $\alpha_k$  are a Robbins-Monroe sequence if:


$$\sum_{k=0:\infty} \alpha_k = \infty$$

$$\sum_{k=0:\infty} \alpha_k^2 < \infty$$

• These conditions are sufficient to ensure convergence of the  $\mathbf{w}_k$  to a **local minimum** of the error function.

E.g.  $\alpha_k = 1 / (k + 1)$  (averaging)

E.g.  $\alpha_k = 1/2$  for  $k = 1, \dots, T$

$\alpha_k = 1/2^2$  for  $k = T+1, \dots, (T+1)+2T$


etc.

---

# Local minima

---

- Convergence is **NOT** to a global minimum, only to local minimum.




- The blue line represents the **error function**. There is no guarantee regarding the amount of error of the weight vector found by gradient descent, compared to the globally optimal solution.

---

# Local minima

---

- Convergence is **NOT** to a global minimum, only to local minimum.




*linear has no local*

- For linear function approximations using Least-Mean Squares (LMS) error, this is not an issue: **only ONE global minimum!**
  - Local minima affects many other function approximators.

# Local minima

- Convergence is **NOT** to a global minimum, only to local minimum.



- For linear function approximations using Least-Mean Squares (LMS) error, this is not an issue: **only ONE global minimum!**
  - Local minima affects many other function approximators.
- Repeated random restarts can help (in all cases of gradient search).

QR

## A 3<sup>rd</sup> optimization method: QR decomposition (optional)

- Consider the usual criteria:
- Assume  $X$  can be decomposed:  
where  $Q$  is an  $n \times m$  orthogonal matrix (i.e.  $Q^T Q = I$ ), and  $R$  is an  $m \times m$  upper triangular matrix.
- Replace  $X$  in equation above:
- Distribute the transpose:
- Let  $Q^T Q = I$  and multiply by  $(R^T)^{-1}$
- Solution:  $\hat{w} = R^{-1} Q^T Y$  The fitted outputs are:  $\hat{Y} = Q Q^T Y$
- This method is more numerically stable than others, and  $R^{-1}$  is fast to compute because upper triangular.
- Alternately, we can use **singular value decomposition**.

$$X^T(Y - Xw) = 0$$

$$X = QR$$



$$(QR)^T Y = (QR)^T (QR) w$$

$$R^T Q^T Y = R^T Q^T Q R w$$

$$Q^T Y = R w$$

---

# What you should know

---

- Definition and characteristics of a supervised learning problem.
- Linear regression (hypothesis class, cost function, algorithm).
- Closed-form least-squares solution method (algorithm, computational complexity, stability issues).
- Gradient descent method (algorithm, properties).

---

# To-do

---

- Reproduce the linear regression example (slides 15-18), solving it using the software of your choice.
- Suggested complementary readings:
  - Ch.2 (Sec. 2.1-2.4, 2.9) of Hastie et al.
  - Ch.3 of Bishop.
  - Ch.9 of Shalev-Schwartz et al.
- Write down **midterm** date in agenda: Nov. 22, 6-8pm, Leacock 132.
- Tutorial times (appearing soon): [www.cs.mcgill.ca/~jpineau/comp551/schedule.html](http://www.cs.mcgill.ca/~jpineau/comp551/schedule.html)
- Office hours (confirmed): [www.cs.mcgill.ca/~jpineau/comp551/syllabus.html](http://www.cs.mcgill.ca/~jpineau/comp551/syllabus.html)